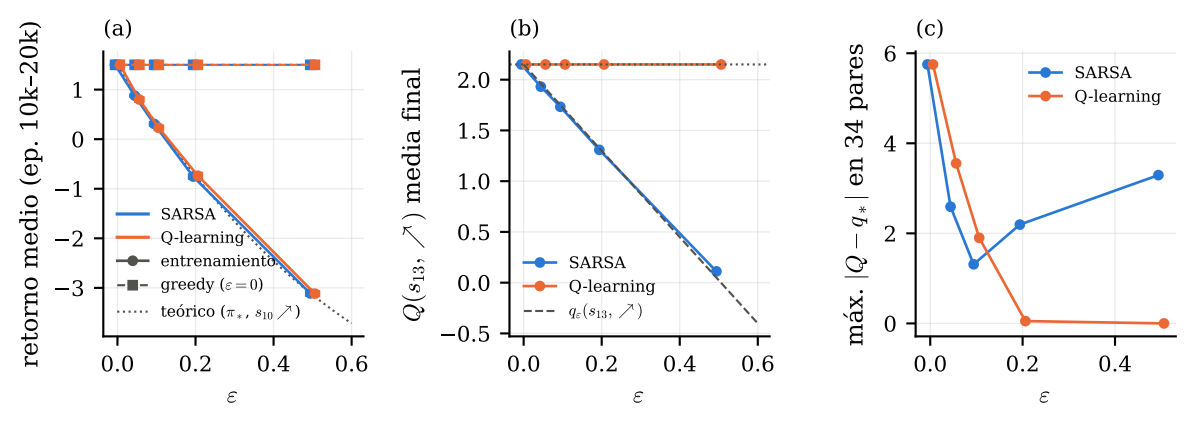
(a)
retorno medio (ep. 10k–20k)
SARSA
Q-learning
entrenamiento
greedy ($\varepsilon = 0$)
teórico (π_*, s_{10} ↗)
ε
(b)
$Q(s_{13}, \nearrow)$ media final
SARSA
Q-learning
$q_\varepsilon(s_{13}, \nearrow)$
ε
(c)
máx. $|Q - q_*|$ en 34 pares
SARSA
Q-learning
ε